Data Quality Framework using PySpark

Project Presentation

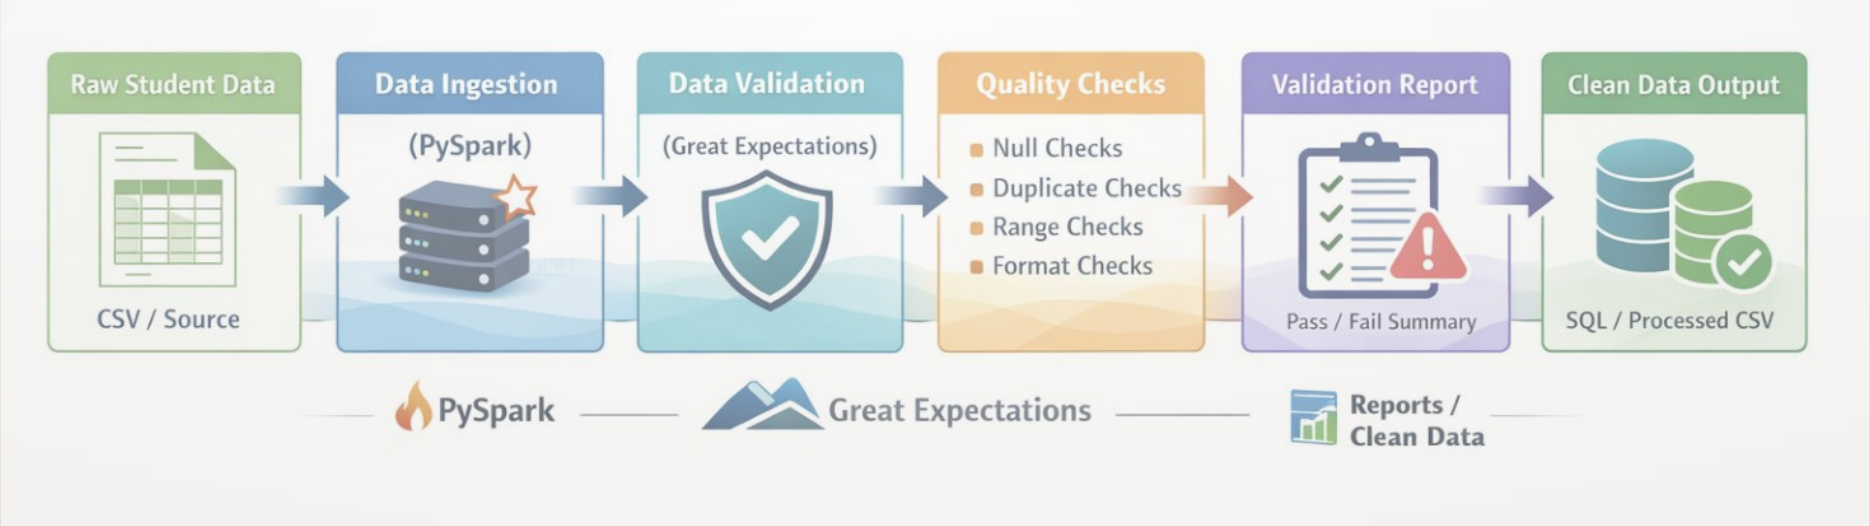

TEAM- 2

Esha-2320030381

Charvi-2320030371

Rushita-2320030064

KL ACCREDITED BY NAAC WITH A++ GRADE

nirf 2022 NATIONAL INSTITUTIONAL RANKING FRAMEWORK RANKED 27 AMONG ALL UNIVERSITIES

CATEGORY 1 UNIVERSITY BY MHRD, Govt. of India

43 YEARS OF EDUCATIONAL LEADERSHIP

CONTENTS

- Abstract
- Introduction
- Literature Survey
- Limitations and Proposed Solutions
- Methodology / Objectives
- Implementation
- Testing
- Deployment

ABSTRACT

Data quality plays a vital role in ensuring accurate analysis, reliable reporting, and effective decision-making in modern data-driven systems. However, real-world datasets often contain various data quality issues such as missing values, duplicate records, invalid formats, and inconsistent or out-of-range values. These issues can negatively impact analytics results, reduce system reliability, and lead to incorrect business insights. Therefore, it is essential to implement mechanisms that can automatically detect and manage data quality problems before data is used for further processing.

This project presents the design and implementation of an automated Data Quality Framework that validates datasets and ensures data integrity. The framework is developed using Python, PySpark, and Great Expectations, which are widely used tools in big data and data engineering environments. It ingests raw data, applies a set of predefined validation rules such as null checks, uniqueness constraints, range validation, and format verification, and identifies records that do not meet quality standards. The framework also generates validation reports that provide detailed information about data quality issues.

The system follows a modular and scalable approach, making it suitable for integration into real-world data pipelines. By automating the data validation process, the framework reduces manual effort, improves efficiency, and ensures that only clean and reliable data is used for analytics and decision-making. This project demonstrates key data engineering concepts such as data ingestion, validation, and reporting, and highlights the importance of maintaining high data quality in building robust and trustworthy data systems.

INTRODUCTION

In modern data-driven systems, the accuracy, completeness, and reliability of data are critical for making correct business decisions and performing meaningful analytics. Poor data quality such as missing values, duplicate records, incorrect formats, or invalid ranges can lead to inaccurate insights, system failures, and loss of trust in data systems. Therefore, ensuring high data quality is an essential responsibility of data engineers.

This project focuses on designing and implementing a Data Quality Framework that automatically validates datasets before they are used for analysis or downstream processing. The framework performs a series of predefined validation checks, such as null value detection, uniqueness constraints, range validation, and format verification, to identify and report data quality issues.

The framework is built using Python, PySpark, and Great Expectations, which are widely used tools in modern data engineering workflows. It ingests raw data, applies data quality rules, generates validation reports, and helps ensure that only clean and reliable data proceeds further in the data pipeline.

This project demonstrates key data engineering concepts, including data ingestion, validation, transformation, and reporting. By implementing this framework, the project highlights the importance of maintaining data integrity and provides a scalable solution that can be integrated into real-world data pipelines to improve data reliability and trustworthiness.

Existing System

In many organizations, data is collected from multiple sources such as databases, logs, APIs, and user inputs. However, these datasets often contain inconsistencies such as duplicate records, missing values, incorrect formats, and invalid ranges. The existing systems usually store and process this data without performing proper validation or quality checks.

Due to the absence of automated data validation mechanisms, poor-quality data enters analytics pipelines and decision-making systems. This results in inaccurate analysis, unreliable insights, and operational inefficiencies. Manual data validation is often time-consuming and prone to human error, making it difficult to maintain data integrity at scale.

Therefore, traditional systems lack automated frameworks to continuously monitor and validate the quality of incoming data.

Limitations of Existing Frameworks and our proposed Solutions

Existing Drawbacks	How our project overcame it
Limited to intrinsic quality only	Add multiple quality dimensions; completeness, accuracy, consistency, validity, timeliness
No real-time validation	Add streaming validation using Spark Structured Streaming
Static validation rules	Implement dynamic rule engine using config file(JSON/YAML)
Limited scalability	Use PySpark distributed processing
No automated reporting	Generate automated quality reports(CSV/HTML)
Domain-specific frameworks	Build domain-independent engine
No software implementation	Full working implementation with modular code
Cannot detect integrated data errors	Add cross-column and cross-table validation
Manual quality checks	Fully automated pipeline

LITERATURE SURVEY

LINK	JOURNAL	DRAWBACKS	FUTURE DIRECTIONS
https://pmc.ncbi.nlm.nih.gov/articles/PMC8019177/	BMC Medical Research Methodology	Limited to intrinsic data quality; excludes task-specific or contextual aspects like currency in EHRs; relies on metadata availability.	Extend to other study types and data quality aspects; integrate more process variables for broader applications.
https://pmc.ncbi.nlm.nih.gov/articles/PMC6568139/	AMIA Joint Summits on Translational	Requires predefined domain ontologies, which demand significant upfront work; does not assess all quality aspects like timeliness or correctness fully; tested on moderate-sized datasets only.	Make HDQF open-source; test on larger national datasets across organizations; develop additional measures for timeliness and correctness.
https://arxiv.org/abs/2403.00526	Science Proceedings arXiv preprint	Systematic assessment of dimensions like accuracy is largely absent; challenges in gauging cleaning success; limited to proposed facets without full empirical validation.	Develop comprehensive framework overcoming data/source/system/task/human facets; align with regulations like EU AI Act.
https://pmc.ncbi.nlm.nih.gov/articles/PMC11297942/	Not specified (PMC)	Focused on ML/DL in medicine; adoption challenges in high-stakes areas due to unproven trustworthiness.	Broaden trustworthy AI development for medical ML applications.
https://pmc.ncbi.nlm.nih.gov/articles/PMC5977591/	Not specified (PMC)	Modified existing framework for specific domain (heart failure); may not generalize without re-operationalization.	Apply to broader clinical domains beyond heart failure.

LITERATURE SURVEY

LINK	JOURNAL	DRAWBACKS	FUTURE DIRECTIONS
https://pmc.ncbi.nlm.nih.gov/articles/PMC10646672/	Not specified (PMC)	Limited to digital health tech interventions; gaps in evaluating real-world health research conduct.	Identify more interventions for data-based health research
https://unstats.un.org/unsd/ccsa/cdqio-2010/Ses1-DQAF-IMF.pdf	IMF Document	Covers prerequisites and five dimensions but generic; lacks software implementations for modern use.	Update for current statistical production needs.
https://nces.ed.gov/fcsm/pdf/FCSM.20.04_A_Framework_for_Data_Quality.pd	FCSM Report	Identifies threats but needs methods for cumulative errors in integrated data; limited disclosure risk assessment.	Develop methods for integrated data errors, gold standards, and disclosure risk
https://gsconlinepress.com/journals/gscarr/content/data-quality-metrics-and-frameworks-comprehensive-study-ensuring-reliable-inf	GSCARR	Persistent challenges in big data accuracy/completeness despite growth; framework alignment with real-time analytics.	Propose integrated DQM for big data, warehousing, real-time analytics
https://medinform.jmir.org/2025/1/e75275	JMIR Medical Informatics	Static evaluations don't adapt to changing data pipelines; limited to core dimensions like completeness.	Advance AIDAVA for dynamic health data integration.

THANK YOU

Team 2